

Machine Learning for Beginners: An Introduction to Neural Networks
A simple explanation of how they work and how to implement one from scratch in Python

by Victor Zhou

https://victorzhou.com/blog/intro-to-neural-networks/?utm_medium=email&utm_source=topic+optin&utm_campaign=awareness&utm_content=20190309+programming+nl&mkt_tok=eyJpIjoiTnpnMk1XUTJNVFU0TURCbSIzInQiOiI0b1FpNHduZ0tnXC9qMXdNa09OVEZkMmJNZmxuN2UxUWtVb0xCQzNcLzVFOU5idVJvY3lIQXlha0FXZ04xN3NtTFNpOVZ0bmRnUmZDb1NEXC9oUnRHeDM0c1MremZlN0NudmJQM1R2OFBoUFBxaGJZUGZhK28xTWVoTlZ6MlpwVUptWiJ9


March 03, 2019

Here's something that might surprise you: neural networks aren't that complicated! The term "neural network" gets used as a buzzword a lot, but in reality they're often much simpler than people imagine.

This post is intended for complete beginners and assumes ZERO prior knowledge of machine learning. We'll understand how neural networks work while implementing one from scratch in Python.

1. Building Blocks: Neurons

First, we have to talk about neurons, the basic unit of a neural network. A neuron takes inputs, does some math with them, and produces one output. Here's what a 2-input neuron looks like:



(see the rest of the material on line, URL is given on the top of this page.)